

# Delivery Seam Review

A lightweight review for one AI delivery slice. Use before build assumptions harden, before release, and at closeout. The purpose is to expose handoff risk and qualify reusable learning—not add ceremony.

|                         |                             |
|-------------------------|-----------------------------|
| Engagement / workstream | Outcome / workflow baseline |
| Technical owner(s)      | Client decision owner       |

## 1. Seam Map

| Seam                   | Question that must be answerable                                                                                  | Evidence / owner / escalation |
|------------------------|-------------------------------------------------------------------------------------------------------------------|-------------------------------|
| Intent ↔ Data          | What outcome and user behavior are changing, and what data boundary makes the claim testable?                     |                               |
| Data ↔ AI              | What retrieval, quality, freshness, metadata, privacy, and access assumptions shape model behavior?               |                               |
| AI ↔ Application       | Which decisions belong to prompts/agents versus deterministic services, UI constraints, or human review?          |                               |
| Application ↔ Platform | What interfaces, deployment, performance, observability, security, and cost constraints can break the experience? |                               |
| System ↔ Human         | Where does human authority remain, what triggers escalation, and how are overrides learned from?                  |                               |
| Delivery ↔ Reuse       | What is reusable, what depends on client context, and what failure modes must travel with the pattern?            |                               |

## 2. Proof Before Release

|                                                                        |                                                                        |
|------------------------------------------------------------------------|------------------------------------------------------------------------|
| <input type="checkbox"/> Golden task cases defined and versioned       | <input type="checkbox"/> Adversarial / boundary cases exercised        |
| <input type="checkbox"/> Groundedness / task-success criteria agreed   | <input type="checkbox"/> Safety and policy evidence retained           |
| <input type="checkbox"/> Privacy and data-minimization review complete | <input type="checkbox"/> Tool permissions and fallbacks bounded        |
| <input type="checkbox"/> Latency and cost assumptions tested           | <input type="checkbox"/> Human review / override / escalation explicit |
| <input type="checkbox"/> Telemetry separates usage from value          | <input type="checkbox"/> Release threshold and rollback owner explicit |

## 3. Closeout Harvest

| Harvest question                                                                                              | Decision |
|---------------------------------------------------------------------------------------------------------------|----------|
| What problem pattern repeated?                                                                                |          |
| What artifact is reusable: agent, accelerator, component, framework, evaluation asset, or cautionary pattern? |          |
| What dependencies and adaptation boundaries must be preserved?                                                |          |
| What evidence from real use supports reuse?                                                                   |          |
| Who owns versioning and feedback?                                                                             |          |

**Candidate thesis:** AI delivery scales at the seams. This worksheet is a hypothesis for discovery and is intended to complement—not replace—3Cloud's existing delivery, quality, architecture, and MILO mechanisms.