AI delivery scales at the seams.

Interview thesis brief for Senior Consultant - AI: mandate, tensions, Delivery Weave, proof map, objection, entry plan, and executive questions.

Russell Dudek Independent candidate brief July 12, 2026

1. Role Intelligence

Nominal role

Senior Consultant - AI supporting architecture and delivery across modernization and transformation programs, with contribution across implementation, cross-practice integration, quality, mentoring, presales, utilization, thought leadership, and reusable IP.

Actual mandate

> **Turn cross-practice architecture into reliable AI delivery, then convert project learning into governed, reusable patterns that improve client value and delivery leverage at post-acquisition scale.**

Company moment

3Cloud became part of Cognizant effective January 1, 2026. Public positioning connects Azure specialization, Data & AI, App Innovation, Cloud Platform, and MILO's agents, accelerators, components, and frameworks. The role therefore carries both engagement-level delivery and system-level learning obligations.

Operating tensions

Tension	Leadership burden
Architecture intent vs delivery reality	Translate guidance into implementable, testable increments without losing coherence.
Client specificity vs reusable IP	Preserve contextual fit while qualifying learning for reuse.
AI speed vs evaluation and trust	Prevent fast iteration from becoming accelerated rework or unmanaged risk.
Autonomy/utilization vs quality/ mentoring	Meet individual delivery economics while improving shared capability.
3Cloud identity vs Cognizant scale	Hypothesis for discovery: scale what helps without diluting delivery behaviors that created value.

Change burden

Teams must make workflow outcomes, data boundaries, human authority, evaluation criteria, telemetry, and reuse boundaries visible early enough to affect delivery—without adding documentation theater.

2. The Delivery Weave

Data & AI, App Innovation, and Cloud Platform pass through five decision gates. Each gate exposes a different integration assumption.

Gate	Core decision	Evidence before moving
Ground	What outcome matters? What data boundary and human authority apply?	Baseline, workflow map, data contract, risk assumptions, evaluation seed set.
Compose	What belongs in retrieval, tools, orchestration, deterministic services, UI, and review?	Interface contracts, bounded permissions, fallback paths, architecture decisions, cost assumptions.
Prove	What is acceptable for task success, groundedness, safety, privacy, latency, and cost?	Golden/adversarial cases, failure taxonomy, human review criteria, release thresholds.
Observe	Which signals distinguish value, adoption, model issues, and workflow/integration issues?	Outcome metrics, adoption depth, overrides, escalations, defects, latency, unit economics.
Harvest	What should become a component, accelerator, framework, evaluation asset, or warning?	Dependencies, adaptation boundaries, failure modes, owner, version, real-use evidence.

Decision architecture

- What remains client-specific and what becomes standardized?
- Which responsibilities belong to AI, deterministic services, or human authority?
- What gets evaluated before release and monitored continuously?
- What cross-practice risk deserves escalation?
- What gets stopped because value, quality, security, or economics do not justify continuation?
- What learning is mature enough for MILO/IP review?

Balanced scorecard

Delivery: architecture-to-accepted increment, predictability, escaped defects, integration rework. **AI quality:** task success, groundedness, safety, review agreement. **Economics:** unit cost, avoided rework, delivery leverage. **Adoption:** workflow penetration, overrides, abandonment. **Trust:** decision traceability, permissions, review coverage. **Reuse:** harvest-to-reuse latency and adaptation burden.

3. Proof Map & Objection

Verified evidence	Transfer to mandate	Claim boundary
Vape-Jet AI-first operating transformation across operations, quality, support, engineering issue flow, ERP/Odoo, Jira, HubSpot, Slack, and telemetry.	Cross-system integration, signal-to-work mechanisms, standard work, adoption, escalation, operating feedback loops.	Direct operating leadership; not Azure consulting delivery.
DudeWorth AI readiness and agentic workflow patterns across context, retrieval, agents, human judgment, escalation, evaluation, and governance.	Solution decomposition, value framing, governance rails, human authority, reuse logic, time-to-value.	Direct framework/advisory work; concepts remain concepts when not deployed.
Amazon 24/7 customer-backward operations supporting approximately five million second-day orders annually.	Cadence, visibility, standard work, escalation, throughput, leadership routines, customer-outcome orientation.	Operations leadership; innovation contributions remain contributed concepts.
Compunetics Engineering, manufacturing, quality, customer delivery, high-reliability technical programs, controls.	Technical translation, design-to-delivery discipline, quality gates, root cause, consequence-aware execution.	Direct advanced-electronics leadership; no conversion into cloud claims.

Strongest plausible objection

> **“This role asks for 2-5 years of Azure delivery. Russell's verified evidence is broader AI and operating transformation, not multi-year Azure AI consulting.”**

Response strategy

Agree with the fact pattern. Do not hide it. Demonstrate transferable disciplines without claiming equivalence: agentic workflow patterns, retrieval, human-in-the-loop design, evaluation and governance rails, cross-system integration, quality systems, technical-business translation, and delivery mechanisms. Ask to be tested through bounded technical casework. Enter with an explicit Azure-specific learning plan and earn scope through demonstrated delivery.

4. Entry Plan, Questions & Sources

Days 0-30 · Learn the house style and earn technical trust

- Map architecture, delivery, quality, review, and MILO contribution paths.
- Pair on active design reviews and close Azure-specific gaps against the real project stack.
- Create one seam-risk map and contribute within existing team norms.

Days 31-60 · Own a bounded workstream

- Translate architecture into tasks, acceptance criteria, interfaces, and risk checks.
- Stand up evaluation and telemetry before release pressure makes them optional.
- Coordinate one cross-practice interface and mentor through pairing and review.

Days 61-90 · Turn delivery learning into leverage

- Separate client context from reusable mechanisms.
- Package one evidence-backed candidate asset for MILO/IP review.
- Share learning and propose the next technical-depth plan.

Executive questions

1. Where do AI engagements most often lose time today?
2. What distinguishes a technically strong Senior Consultant from one who improves the delivery system?
3. How does MILO qualify reusable IP?
4. Where should 3Cloud identity remain distinctive after the acquisition?
5. What evidence earns trust across the three practices?
6. What would prove after six months that this person reduced risk and improved others?

Public sources

3Cloud job posting; 3Cloud acquisition announcement; MILO overview; Data Science and AI services; About 3Cloud. Internal recommendations remain hypotheses for discovery.